

Midterm**Index**

1. Deployment Plan	2
2. Accessing Instances through SSH	3
3. MongoDB Installation	3
4. Data Directories	4
5. Private & Public IPs of Nodes	4
6. Config Server Set-up	5
7. Connect to Mongos	7
8. Shard Set-up	7
9. Adding Shards	10
10. Shard Key & Sharding Strategy	11
11. Adding Data from Public Dataset	11
12. Running Queries on Data	13
1. Find documents in a given range	13
2. Query with \$elemMatch with two conditions	15
3. Query with \$in, \$nin, or \$all	17
4. Query involving aggregate()	19
5. Query involving mapReduce()	19
6. Update query	20
13. Replication of Shards	20
14. List of Hosts	24

[Link to Doc/PDF](#)

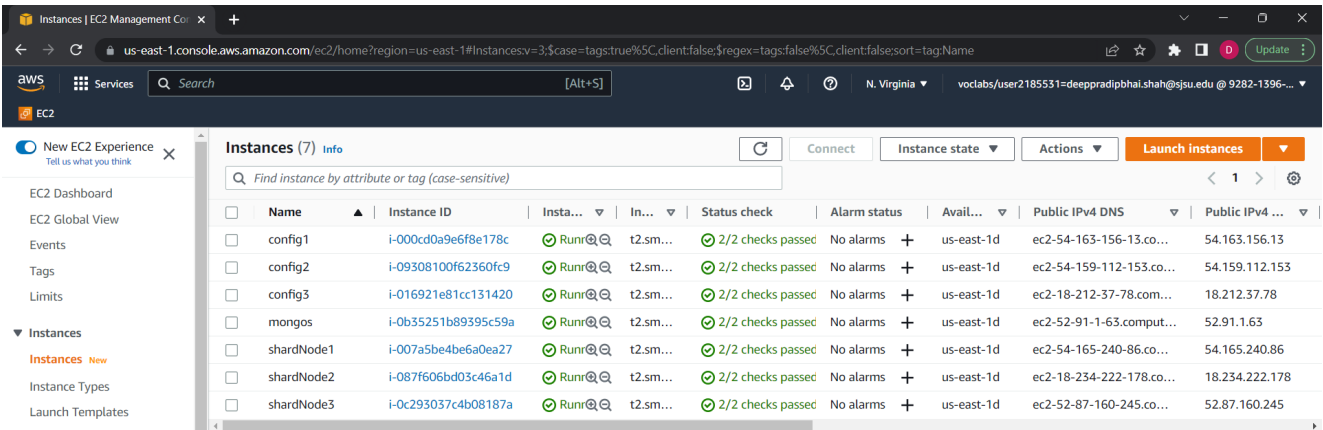
1. Deployment Plan

Number of Nodes: 7

Deployment Summary

Service	Instance/ Node Name	IP Address (Private)	Port	Replica Set Name	Daemon	DB Path	Replica Priority
Shard Controller	mongos	172.31.80.79	27020	-	mongos	-	-
Config Server Primary	config1	172.31.90.97	27019	crs	mongod	/data/db	1
Config Server Secondary 1	config2	172.31.89.25	27019	crs	mongod	/data/db	1
Config Server Secondary 2	config3	172.31.86.144	27019	crs	mongod	/data/db	1
Shard a Primary	shardNode1	172.31.93.248	27021	a	mongod	/data/a	4
Shard b Secondary 1		172.31.93.248	27022	b	mongod	/data/b	1
Shard c Secondary 1		172.31.93.248	27023	c	mongod	/data/c	1
Shard a Secondary 1	shardNode2	172.31.84.139	27021	a	mongod	/data/a	1
Shard b Primary		172.31.84.139	27022	b	mongod	/data/b	4
Shard c Secondary 2		172.31.84.139	27023	c	mongod	/data/c	1
Shard a Secondary 2	shardNode3	172.31.90.164	27021	a	mongod	/data/a	1
Shard b Secondary 2		172.31.90.164	27022	b	mongod	/data/b	1
Shard c Primary		172.31.90.164	27023	c	mongod	/data/c	4

Nodes Set-up in AWS:

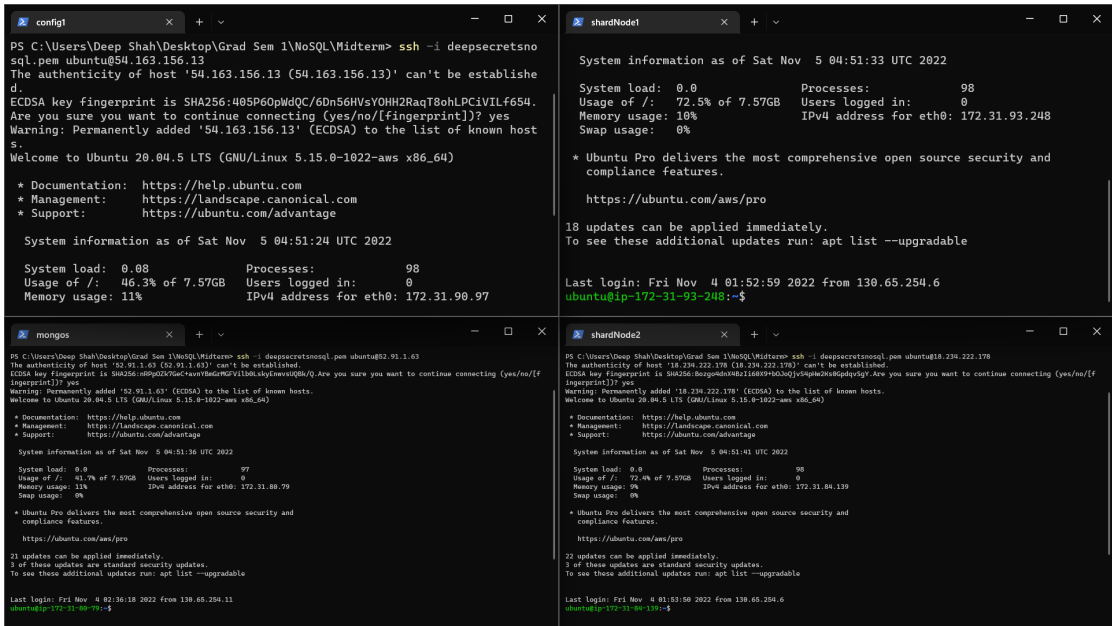


2. Accessing Instances through SSH

Command:

```
ssh -i secret_key.pem ubuntu@{public_ip}
```

Accessing some instances through SSH:

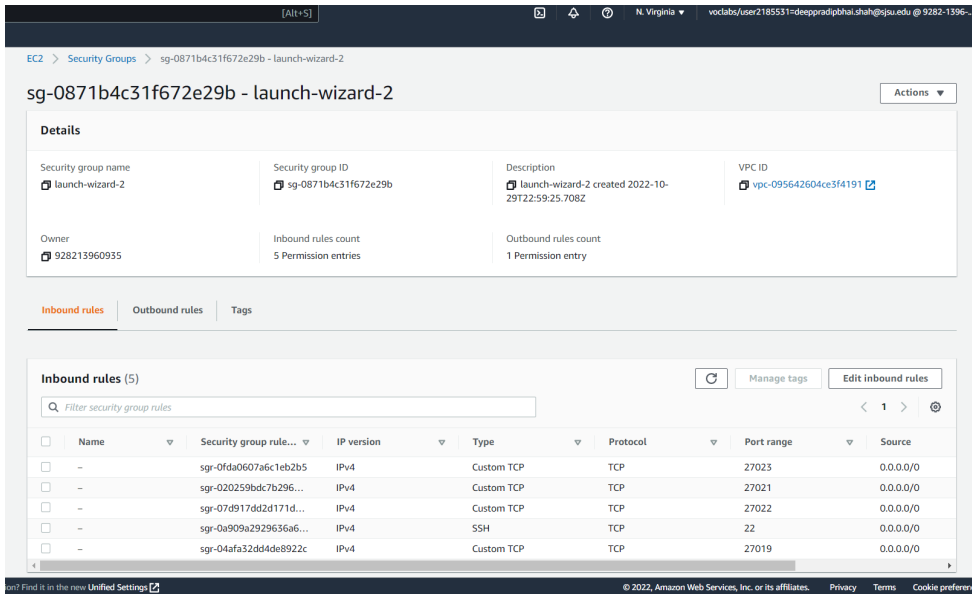


3. MongoDB Installation

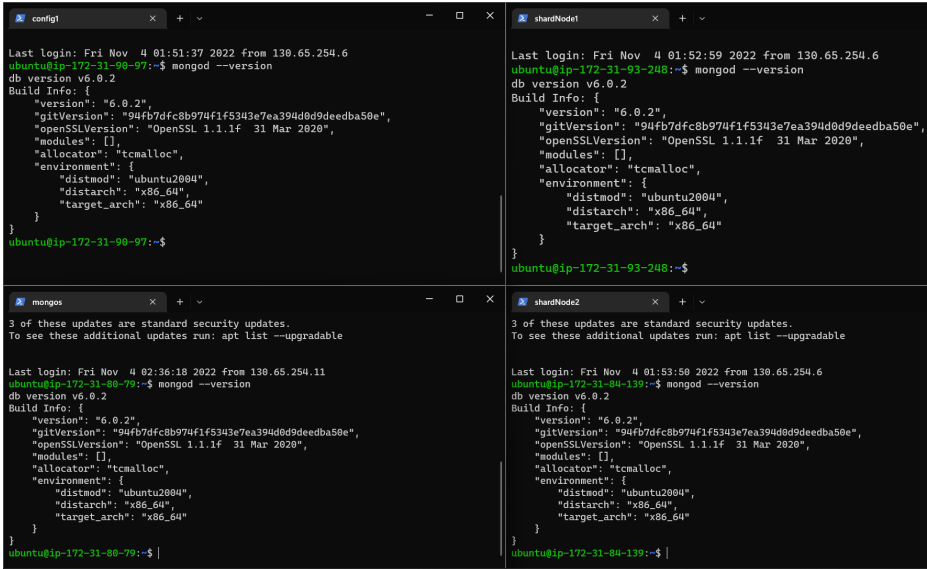
Step Followed: [Install MongoDB on Ubuntu](#)

Also, enable the inbound network in the security group of each node.

Allowing Inbound Network traffic on certain ports in the Security Group:



Checking the mongod version to verify MongoDB Installation:



The screenshot displays four terminal windows arranged in a 2x2 grid. The top-left window, titled 'config1', shows the command 'mongod --version' being executed on node 'config1' (IP 172.31.90.97), returning 'db version v6.0.2' and detailed build information. The top-right window, titled 'shardNode1', shows the same command on node 'shardNode1' (IP 172.31.93-248), also returning 'db version v6.0.2' and build info. The bottom-left window, titled 'mongos', shows the command 'mongod --version' on node 'mongos' (IP 172.31-80-79), returning 'db version v6.0.2' and build info. The bottom-right window, titled 'shardNode2', shows the command 'mongod --version' on node 'shardNode2' (IP 172-31-84-139), returning 'db version v6.0.2' and build info. All nodes report version 6.0.2.

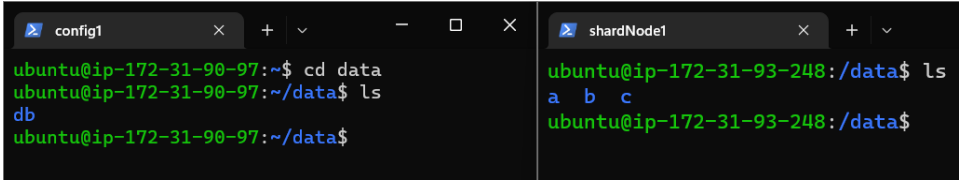
4. Data Directories

Command:

sudo mkdir -p /data/db

sudo mkdir -p /data/a /data/b /data/c

Checking the data directories created:



The screenshot shows two terminal windows. The left window, titled 'config1', shows the commands 'cd data' and 'ls' being executed on node 'config1' (IP 172-31-90-97), resulting in the output 'a b c db'. The right window, titled 'shardNode1', shows the command 'ls' being executed on node 'shardNode1' (IP 172-31-93-248), resulting in the output 'a b c'. Both windows confirm the successful creation of the data directories.

5. Private & Public IPs of Nodes

Instance Name	Private IP Address	Public IP Address
mongos	172.31.80.79	52.91.1.63
config1	172.31.90.97	54.163.156.13
config2	172.31.89.25	54.159.112.153
config3	172.31.86.144	18.212.37.78
shardNode1	172.31.93.248	54.165.240.86
shardNode2	172.31.84.139	18.234.222.178
shardNode3	172.31.90.164	52.87.160.245

6. Config Server Set-up

Commands:

Setup Config servers (to be run in each node):

```
sudo mongod --configsvr --replSet crs --dbpath /data/db --port 27019 --logpath  
/var/log/mongodb/mongod.log --bind_ip_all --fork
```

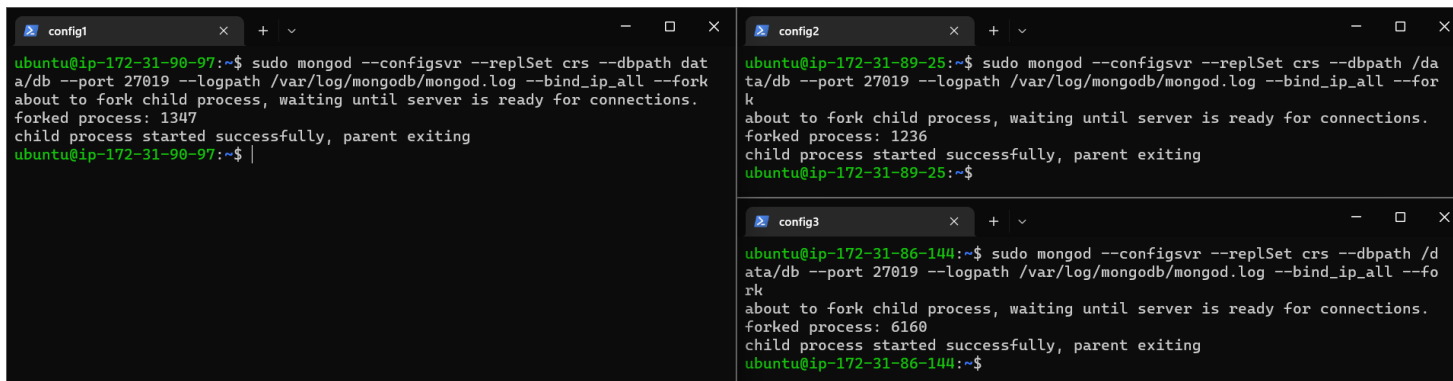
Initiate Replica Set:

```
rs.initiate( { _id: "crs", configsvr: true, members : [ { _id: 0, host: "172.31.90.97:27019"}, { _id: 1, host:  
"172.31.89.25:27019"}, { _id: 2, host: "172.31.86.144:27019"} ] } )
```

Steps:

1. In each of the instances config1, config2, config3, run the above command to Setup the Config Servers.
2. Use command: mongosh -port 27019 in any one instance.
3. Initiate the replica set using the rs.initiate command above.

Running the Config Servers:



```
config1
ubuntu@ip-172-31-90-97:~$ sudo mongod --configsvr --replSet crs --dbpath data/db --port 27019 --logpath /var/log/mongodb/mongod.log --bind_ip_all --fork
about to fork child process, waiting until server is ready for connections.
forked process: 1347
child process started successfully, parent exiting
ubuntu@ip-172-31-90-97:~$

config2
ubuntu@ip-172-31-89-25:~$ sudo mongod --configsvr --replSet crs --dbpath /data/db --port 27019 --logpath /var/log/mongodb/mongod.log --bind_ip_all --fork
about to fork child process, waiting until server is ready for connections.
forked process: 1236
child process started successfully, parent exiting
ubuntu@ip-172-31-89-25:~$

config3
ubuntu@ip-172-31-86-144:~$ sudo mongod --configsvr --replSet crs --dbpath /data/db --port 27019 --logpath /var/log/mongodb/mongod.log --bind_ip_all --fork
about to fork child process, waiting until server is ready for connections.
forked process: 6160
child process started successfully, parent exiting
ubuntu@ip-172-31-86-144:~$
```

Config Server Replica Set:

```

config1 x + - □ ×

crs [direct: secondary] test> rs.status()
{
  set: 'crs',
  date: ISODate("2022-11-05T05:51:51.063Z"),
  myState: 2,
  term: Long("13"),
  syncSourceHost: '172.31.89.25:27019',
  syncSourceId: 1,
  configsvr: true,
  heartbeatIntervalMillis: Long("2000"),
  majorityVoteCount: 2,
  writeMajorityCount: 2,
  votingMembersCount: 3,
  writableVotingMembersCount: 3,
  optimes: {
    lastCommittedOpTime: { ts: Timestamp({ t: 1667627510, i: 1 }), t: Long("13") },
    lastCommittedWallTime: ISODate("2022-11-05T05:51:50.055Z"),
    readConcernMajorityOpTime: { ts: Timestamp({ t: 1667627510, i: 1 }), t: Long("13") }
  },
  appliedOpTime: { ts: Timestamp({ t: 1667627511, i: 1 }), t: Long("13") },
  durableOpTime: { ts: Timestamp({ t: 1667627511, i: 1 }), t: Long("13") },
  lastAppliedWallTime: ISODate("2022-11-05T05:51:51.055Z"),
  lastDurableWallTime: ISODate("2022-11-05T05:51:51.055Z"),
  lastStableRecoveryTimestamp: Timestamp({ t: 1667627496, i: 1 }),
  electionParticipantMetrics: {
    votedForCandidate: true,
    electionTerm: Long("13"),
    lastVoteDate: ISODate("2022-11-05T05:48:51.971Z"),
    electionCandidateMemberId: 1,
    voteReason: ""
  },
  lastAppliedOpTimeAtElection: { ts: Timestamp({ t: 1667540936, i: 1 }), t: Long("11") },
  maxAppliedOpTimeInSet: { ts: Timestamp({ t: 1667540936, i: 1 }), t: Long("11") },
  priorityAtElection: 1,
  newTermStartDate: ISODate("2022-11-05T05:48:51.992Z"),
  newTermAppliedDate: ISODate("2022-11-05T05:48:52.034Z"),
  members: [
    {
      _id: 0,
      name: '172.31.90.97:27019',
      health: 1,
      state: 2,
      stateStr: 'SECONDARY',
      uptime: 196,
      optime: { ts: Timestamp({ t: 1667627511, i: 1 }), t: Long("13") },
      optimeDate: ISODate("2022-11-05T05:51:51.000Z"),
      lastAppliedWallTime: ISODate("2022-11-05T05:51:51.055Z"),
      lastDurableWallTime: ISODate("2022-11-05T05:51:51.055Z"),
      syncSourceHost: '172.31.89.25:27019',
      syncSourceId: 1,
      infoMessage: "",
      configVersion: 1,
      configTerm: 13,
      self: true,
      lastHeartbeatMessage: ""
    },
    {
      _id: 1,
      name: '172.31.89.25:27019',
      health: 1,
      state: 1,
      stateStr: 'PRIMARY',
      uptime: 189,
      optime: { ts: Timestamp({ t: 1667627510, i: 1 }), t: Long("13") },
      optimeDurable: { ts: Timestamp({ t: 1667627510, i: 1 }), t: Long("13") },
      optimeDate: ISODate("2022-11-05T05:51:50.000Z"),
      optimeDurableDate: ISODate("2022-11-05T05:51:50.000Z"),
      lastAppliedWallTime: ISODate("2022-11-05T05:51:50.055Z"),
      lastDurableWallTime: ISODate("2022-11-05T05:51:50.055Z"),
      lastHeartbeat: ISODate("2022-11-05T05:51:50.565Z"),
      lastHeartbeatRecv: ISODate("2022-11-05T05:51:50.045Z"),
      pingMs: Long("0"),
      lastHeartbeatMessage: "",
      syncSourceHost: "",
      syncSourceId: -1,
      infoMessage: "",
      electionTime: Timestamp({ t: 1667627331, i: 1 }),
      electionDate: ISODate("2022-11-05T05:48:51.000Z"),
      configVersion: 1,
      configTerm: 13
    },
    {
      _id: 2,
      name: '172.31.86.144:27019',
      health: 1,
      state: 2,
      stateStr: 'SECONDARY',
      uptime: 183,
      optime: { ts: Timestamp({ t: 1667627510, i: 1 }), t: Long("13") },
      optimeDurable: { ts: Timestamp({ t: 1667627510, i: 1 }), t: Long("13") },
      optimeDate: ISODate("2022-11-05T05:51:50.000Z"),
      optimeDurableDate: ISODate("2022-11-05T05:51:50.000Z"),
      lastAppliedWallTime: ISODate("2022-11-05T05:51:50.055Z"),
      lastDurableWallTime: ISODate("2022-11-05T05:51:50.055Z"),
      lastHeartbeat: ISODate("2022-11-05T05:51:50.565Z"),
      lastHeartbeatRecv: ISODate("2022-11-05T05:51:50.564Z"),
      pingMs: Long("0"),
      lastHeartbeatMessage: "",
      syncSourceHost: '172.31.89.25:27019',
      syncSourceId: 1,
      infoMessage: "",
      configVersion: 1,
      configTerm: 13
    }
  ],
  ok: 1,
  lastCommittedOpTime: Timestamp({ t: 1667627510, i: 1 }),
  '$clusterTime': {
    clusterTime: Timestamp({ t: 1667627511, i: 1 }),
    signature: {
      hash: Binary(Buffer.from("00000000000000000000000000000000", "hex"), 0),
      keyId: Long("0")
    }
  },
  operationTime: Timestamp({ t: 1667627511, i: 1 })
}
crs [direct: secondary] test>

```

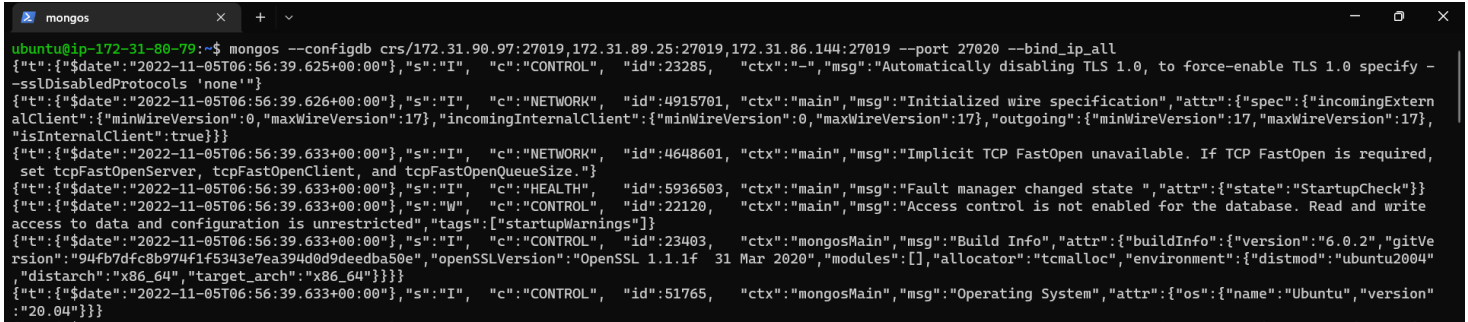
7. Connect to Mongos

Command:

```
mongos --configdb crs/172.31.90.97:27019,172.31.89.25:27019,172.31.86.144:27019 --port 27020 --bind_ip_all
```

Steps: In mongos instance, run the above command.

Connecting Mongos to each config server:



```
ubuntu@ip-172-31-80-79:~$ mongos --configdb crs/172.31.90.97:27019,172.31.89.25:27019,172.31.86.144:27019 --port 27020 --bind_ip_all
{"t":{"sdate":"2022-11-05T06:56:39.625+00:00"},"s":"I", "c":"CONTROL", "id":23285, "ctx":"-", "msg":"Automatically disabling TLS 1.0, to force-enable TLS 1.0 specify --sslDisabledProtocols 'none'"}
{"t":{"sdate":"2022-11-05T06:56:39.626+00:00"},"s":"I", "c":"NETWORK", "id":4915701, "ctx":"main", "msg":"Initialized wire specification", "attr":{"spec":{"incomingExternalClient":{"minWireVersion":0,"maxWireVersion":17},"incomingInternalClient":{"minWireVersion":0,"maxWireVersion":17},"outgoing":{"minWireVersion":17,"maxWireVersion":17},"isInternalClient":true}}}
{"t":{"sdate":"2022-11-05T06:56:39.633+00:00"},"s":"I", "c":"NETWORK", "id":4648601, "ctx":"main", "msg":"Implicit TCP FastOpen unavailable. If TCP FastOpen is required, set tcpFastOpenServer, tcpFastOpenClient, and tcpFastOpenQueueSize."}
{"t":{"sdate":"2022-11-05T06:56:39.633+00:00"},"s":"I", "c":"HEALTH", "id":5936503, "ctx":"main", "msg":"Fault manager changed state ", "attr":{"state":"StartupCheck"}}
{"t":{"sdate":"2022-11-05T06:56:39.633+00:00"},"s":"W", "c":"CONTROL", "id":22120, "ctx":"main", "msg":"Access control is not enabled for the database. Read and write access to data and configuration is unrestricted", "tags":{"startupWarnings":[]}}
{"t":{"sdate":"2022-11-05T06:56:39.633+00:00"},"s":"I", "c":"CONTROL", "id":23403, "ctx":"mongosMain", "msg":"Build Info", "attr":{"buildInfo":{"version":"6.0.2","gitVersion":"94fb7dfc8b974f1f5343e7ea394d0d9deedba50e","openSSLVersion":"OpenSSL 1.1.1f 31 Mar 2020","modules":[],"allocator":"tcmalloc","environment":{"distmod":"ubuntu2004","distarch":"x86_64","target_arch":"x86_64"}}}}
{"t":{"sdate":"2022-11-05T06:56:39.633+00:00"},"s":"I", "c":"CONTROL", "id":51765, "ctx":"mongosMain", "msg":"Operating System", "attr":{"os":{"name":"Ubuntu","version":"20.04"}}}
```

8. Shard Set-up

Shard Set-up: There are 3 shards *a*, *b*, and *c* in each of the nodes shardNode1, shardNode2, and shardNode3.

They are deployed in a replica set; hence, each node has a replica for each shard.

Primary for *a* is shardNode1. Also, the replica set of *a* has priority 4 for shardNode1, and priority 1 for others.

Primary for *b* is shardNode2. Also, the replica set of *b* has priority 4 for shardNode2, and priority 1 for others.

Primary for *c* is shardNode3. Also, the replica set of *c* has priority 4 for shardNode3, and priority 1 for others.

Commands:

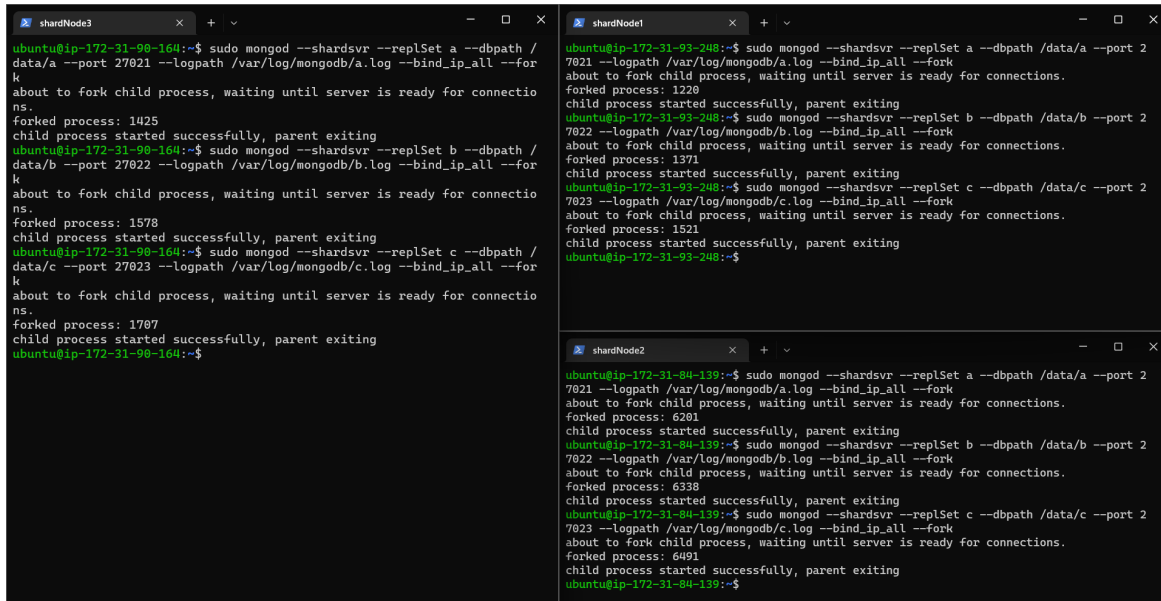
```
sudo mongod --shardsvr --replSet a --dbpath /data/a --port 27021 --logpath /var/log/mongodb/a.log --bind_ip_all --fork
```

```
sudo mongod --shardsvr --replSet b --dbpath /data/b --port 27022 --logpath /var/log/mongodb/b.log --bind_ip_all --fork
```

```
sudo mongod --shardsvr --replSet c --dbpath /data/c --port 27023 --logpath /var/log/mongodb/c.log --bind_ip_all --fork
```

Run the above commands in each of the nodes shardNode1, shardNode2 and shardNode3 to run the shard servers.

Launching 3 shards in a replica set:



```
shardNode3
ubuntu@ip-172-31-90-164:~$ sudo mongod --shardsvr --replSet a --dbpath /data/a --port 27021 --logpath /var/log/mongodb/a.log --bind_ip_all --fork
about to fork child process, waiting until server is ready for connections.
forked process: 1425
child process started successfully, parent exiting
ubuntu@ip-172-31-90-164:~$ sudo mongod --shardsvr --replSet b --dbpath /data/b --port 27022 --logpath /var/log/mongodb/b.log --bind_ip_all --fork
about to fork child process, waiting until server is ready for connections.
forked process: 1578
child process started successfully, parent exiting
ubuntu@ip-172-31-90-164:~$ sudo mongod --shardsvr --replSet c --dbpath /data/c --port 27023 --logpath /var/log/mongodb/c.log --bind_ip_all --fork
about to fork child process, waiting until server is ready for connections.
forked process: 1707
child process started successfully, parent exiting
ubuntu@ip-172-31-90-164:~$

shardNode1
ubuntu@ip-172-31-93-248:~$ sudo mongod --shardsvr --replSet a --dbpath /data/a --port 27021 --logpath /var/log/mongodb/a.log --bind_ip_all --fork
about to fork child process, waiting until server is ready for connections.
forked process: 1220
child process started successfully, parent exiting
ubuntu@ip-172-31-93-248:~$ sudo mongod --shardsvr --replSet b --dbpath /data/b --port 27022 --logpath /var/log/mongodb/b.log --bind_ip_all --fork
about to fork child process, waiting until server is ready for connections.
forked process: 1371
child process started successfully, parent exiting
ubuntu@ip-172-31-93-248:~$ sudo mongod --shardsvr --replSet c --dbpath /data/c --port 27023 --logpath /var/log/mongodb/c.log --bind_ip_all --fork
about to fork child process, waiting until server is ready for connections.
forked process: 1521
child process started successfully, parent exiting
ubuntu@ip-172-31-93-248:~$

shardNode2
ubuntu@ip-172-31-84-139:~$ sudo mongod --shardsvr --replSet a --dbpath /data/a --port 27021 --logpath /var/log/mongodb/a.log --bind_ip_all --fork
about to fork child process, waiting until server is ready for connections.
forked process: 6201
child process started successfully, parent exiting
ubuntu@ip-172-31-84-139:~$ sudo mongod --shardsvr --replSet b --dbpath /data/b --port 27022 --logpath /var/log/mongodb/b.log --bind_ip_all --fork
about to fork child process, waiting until server is ready for connections.
forked process: 6338
child process started successfully, parent exiting
ubuntu@ip-172-31-84-139:~$ sudo mongod --shardsvr --replSet c --dbpath /data/c --port 27023 --logpath /var/log/mongodb/c.log --bind_ip_all --fork
about to fork child process, waiting until server is ready for connections.
forked process: 6491
child process started successfully, parent exiting
ubuntu@ip-172-31-84-139:~$
```

To initiate replica set for shard *a*, run commands:

```
mongosh -port 27021
```

```
rs.initiate({_id: "a", members : [ { _id: 0, host: "172.31.93.248:27021"}, { _id: 1, host: "172.31.84.139:27021"}, { _id: 2, host: "172.31.90.164:27021"} ] })
```

Change Priority of replica set member:

```
conf = rs.conf()
```

```
conf.members[0].priority = 4
```

```
rs.reconfig(conf)
```

To initiate replica set for shard *b*, run commands:

```
mongosh -port 27022
```

```
rs.initiate({_id: "b", members : [ { _id: 0, host: "172.31.93.248:27022"}, { _id: 1, host: "172.31.84.139:27022"}, { _id: 2, host: "172.31.90.164:27022"} ] })
```

Change Priority of replica set member:

```
conf = rs.conf()
```

```
conf.members[1].priority = 4
```

```
rs.reconfig(conf)
```


To initiate replica set for shard *c*, run commands:

```
mongosh -port 27023
```

```
rs.initiate({_id: "c", members : [ {_id: 0, host: "172.31.93.248:27023"}, {_id: 1, host: "172.31.84.139:27023"}, {_id: 2, host: "172.31.90.164:27023"} ] } )
```

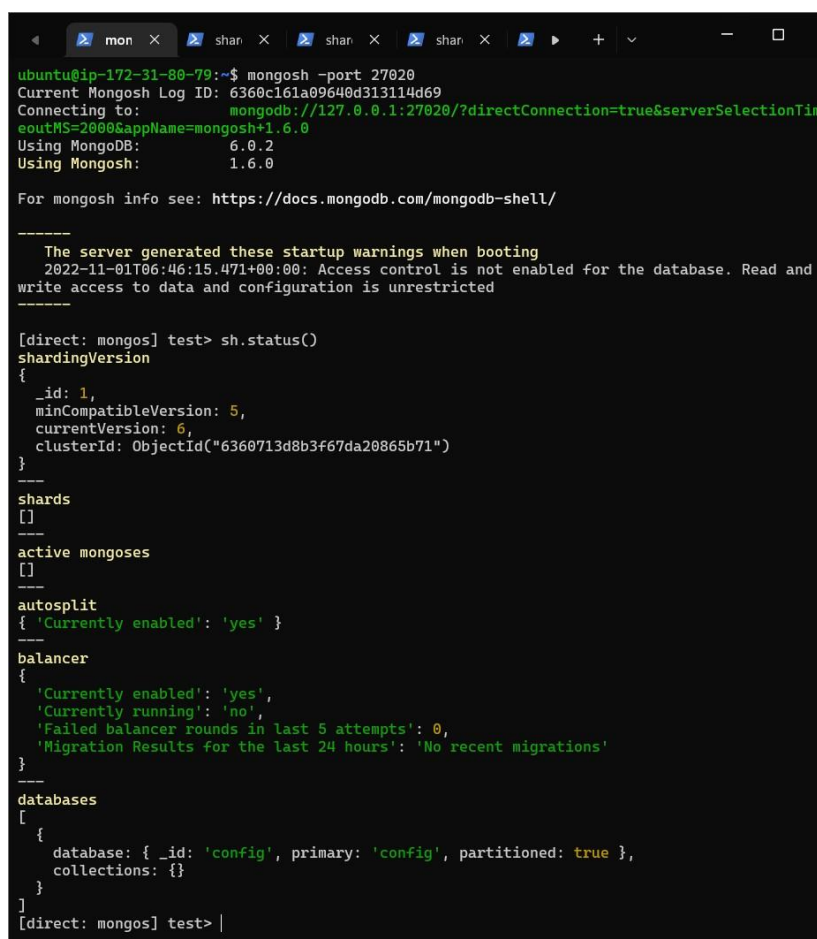
Change Priority of replica set member:

```
conf = rs.conf()
conf.members[2].priority = 4
rs.reconfig(conf)
```

I changed the priority of the replica set to ensure that the primary of every node is on a different shard and one node is handling requests from only one shard when all nodes are healthy. Also, I kept replicas on different nodes to ensure the availability of all shards in case of failure.

Shard status from mongos:

(Before adding shards)



```
ubuntu@ip-172-31-80-79:~$ mongosh -port 27020
Current Mongosh Log ID: 6360c161a09640d313114d69
Connecting to:      mongodb://127.0.0.1:27020/?directConnection=true&serverSelectionTim
eoutMS=2000&appName=mongosh+1.6.0
Using MongoDB:      6.0.2
Using Mongosh:      1.6.0

For mongosh info see: https://docs.mongodb.com/mongosh-shell/

-----
The server generated these startup warnings when booting
2022-11-01T06:46:15.471+00:00: Access control is not enabled for the database. Read and
write access to data and configuration is unrestricted
-----

[direct: mongos] test> sh.status()
shardingVersion
{
  _id: 1,
  minCompatibleVersion: 5,
  currentVersion: 6,
  clusterId: ObjectId("6360713d8b3f67da20865b71")
}
-----
shards
[]
-----
active mongoses
[]
-----
autosplit
{ 'Currently enabled': 'yes' }
-----
balancer
{
  'Currently enabled': 'yes',
  'Currently running': 'no',
  'Failed balancer rounds in last 5 attempts': 0,
  'Migration Results for the last 24 hours': 'No recent migrations'
}
-----
databases
[
  {
    database: { _id: 'config', primary: 'config', partitioned: true },
    collections: {}
  }
]
[direct: mongos] test> |
```

9. Adding Shards

In mongos instance, run commands:

```
mongosh -port 27020
```

```
sh.addShard( "a/172.31.93.248:27021,172.31.84.139:27021,172.31.90.164:27021" )
```

```
sh.addShard( "b/172.31.93.248:27022,172.31.84.139:27022,172.31.90.164:27022" )
```

```
sh.addShard( "c/172.31.93.248:27023,172.31.84.139:27023,172.31.90.164:27023" )
```

Shard status After adding shards:

```
ubuntu@ip-172-31-80-79:~$ mongosh -port 27020
Current Mongosh Log ID: 6360c7f0967d65c120c438f0
Connecting to:  mongodb://127.0.0.1:27020/?directConnection=true&serverSelectionTimeoutMS=2000&appName=mongosh+1.6.0
Using MongoDB:  6.0.2
Using Mongosh:  1.6.0

For mongosh info see: https://docs.mongodb.com/mongosh-shell/

-----
The server generated these startup warnings when booting
2022-11-01T06:46:15.471+00:00: Access control is not enabled for the database. Read and write access to data and configuration is unrestricted
-----

[direct: mongos] test> sh.status()
shardingVersion
{
  _id: 1,
  minCompatibleVersion: 5,
  currentVersion: 6,
  clusterId: ObjectId("6360713d8b3f67da20865b71")
}
---
shards
[
  {
    _id: 'a',
    host: 'a/172.31.84.139:27021,172.31.90.164:27021,172.31.93.248:27021',
    state: 1,
    topologyTime: Timestamp({ t: 1667285858, i: 1 })
  },
  {
    _id: 'b',
    host: 'b/172.31.84.139:27022,172.31.90.164:27022,172.31.93.248:27022',
    state: 1,
    topologyTime: Timestamp({ t: 1667285956, i: 2 })
  },
  {
    _id: 'c',
    host: 'c/172.31.84.139:27023,172.31.90.164:27023,172.31.93.248:27023',
    state: 1,
    topologyTime: Timestamp({ t: 1667285967, i: 7 })
  }
]
```

```
state: 1,
topologyTime: Timestamp({ t: 1667285967, i: 7 })
}
]
---
active mongoses
[ { '6.0.2': 1 } ]
---
autosplit
{ 'Currently enabled': 'yes' }
---
balancer
{
  'Currently enabled': 'yes',
  'Currently running': 'no',
  'Failed balancer rounds in last 5 attempts': 0,
  'Migration Results for the last 24 hours': { '342': 'Success' }
}
---
databases
[
  {
    database: { _id: 'config', primary: 'config', partitioned: true },
    collections: {
      'config.system.sessions': {
        shardKey: { _id: 1 },
        unique: false,
        balancing: true,
        chunkMetadata: [
          { shard: 'a', nChunks: 682 },
          { shard: 'b', nChunks: 171 },
          { shard: 'c', nChunks: 171 }
        ],
        chunks: [
          'too many chunks to print, use verbose if you want to force print'
        ],
        tags: []
      }
    }
  }
]
[direct: mongos] test> |
```

10. Shard Key & Sharding Strategy

Shard Key: Random (Field: 'title', index: hashed)

Sharding Strategy: Hash-based

Commands:

```
use nosql_mid
sh.enableSharding("nosql_mid")
db.books.createIndex({"title":"hashed"})
sh.shardCollection("nosql_mid.books", {title:"hashed"})
use config
db.settings.updateOne(
  { _id: "chunksize" },
  { $set: { _id: "chunksize", value: 1 } },
  { upsert: true }
)
```

11. Adding Data from Public Dataset

Dataset URL: <https://www.kaggle.com/datasets/opalskies/large-books-metadata-dataset-50-mill-entries>

File from above URL: list.json

Collection Description:

The collection is about suggestions on books. Every document in this collection has fields like the number of voters, number of likes, an array of descriptions of suggested books, a list of tags, and so on.

One Random Document from the above Dataset:

```
{
  "_id": 29,
  "title": "Best Books of 2007",
  "description": "The best books published during 2007. See best-of-year lists for other ...",
  "description_html": "The best books published during 2007.<br><br>See best-of-year ...",
  "num_pages": 10,
  "num_books": 916,
  "num_voters": 1178,
  "created_date": "June 21st, 2008",
  "tags": Array
    0: "2007"
    1: "best"
  "num_likes": 30,
  "created_by": Object
    name: "deleted user"
    id: ""
  "num_comments": 6,
  "books": Array
    0: Object
      book_id: "136251"
      title: "Harry Potter and the Deathly Hallows"
      author_id: "1077326"
      author: "J.K. Rowling"
      position: Object
        ranking: 1
        score: 37269
        votes: 376
    1: Object
    2: Object
    ...
}
```

To insert data:

1. Download the list.json file from the link above and convert it to json format if needed.
2. Connect to mongos server from the MongoDB-Compass desktop application.
3. Import data in the collection 'books'.

Shard status After populating Data in a sharded collection:

```

mongos
[direct: mongos] nosql_mid> sh.status()
shardingVersion
{
  _id: 1,
  minCompatibleVersion: 5,
  currentVersion: 6,
  clusterId: ObjectId("6360713d8b3f67da20865b71")
}
---
shards
[
  {
    _id: 'a',
    host: 'a/172.31.84.139:27021,172.31.90.164:27021,172.31.93.248:27021',
    state: 1,
    topologyTime: Timestamp({ t: 1667285858, i: 1 })
  },
  {
    _id: 'b',
    host: 'b/172.31.84.139:27022,172.31.90.164:27022,172.31.93.248:27022',
    state: 1,
    topologyTime: Timestamp({ t: 1667285956, i: 2 })
  },
  {
    _id: 'c',
    host: 'c/172.31.84.139:27023,172.31.90.164:27023,172.31.93.248:27023',
    state: 1,
    topologyTime: Timestamp({ t: 1667285967, i: 7 })
  }
]
---
active mongoses
[ { '6.0.2': 1 } ]
---
autosplit
{ 'Currently enabled': 'yes' }
---
balancer
{
  'Currently enabled': 'yes',
  'Currently running': 'no',
  'Failed balancer rounds in last 5 attempts': 0,
  'Migration Results for the last 24 hours': 'No recent migrations'
}
---
databases
[
  {
    database: {
      _id: 'books',
      primary: 'a',
      partitioned: false,
      version: {
        uuid: new UUID("b7f30595-c31e-4572-b9f8-6e6dcccfb8f9"),
        timestamp: Timestamp({ t: 1667363566, i: 7 }),
        lastMod: 1
      }
    },
    collections: {}
  },
  {
    database: { _id: 'config', primary: 'config', partitioned: true },
    collections: {
      'config.system.sessions': {
        shardKey: { _id: 1 },
        unique: false,
        balancing: true,
        chunkMetadata: [
          { shard: 'a', nChunks: 342 },
          { shard: 'b', nChunks: 341 },
          { shard: 'c', nChunks: 341 }
        ],
        chunks: [
          'too many chunks to print, use verbose if you want to force print'
        ],
        tags: []
      }
    }
  },
  {
    database: {
      _id: 'nosql_mid',
      primary: 'a',
      partitioned: false,
      version: {
        uuid: new UUID("8b024bad-aa49-4039-9d88-45db841abd4b"),
        timestamp: Timestamp({ t: 1667363694, i: 3 }),
        lastMod: 1
      }
    },
    collections: {
      'nosql_mid.books': {
        shardKey: { title: 'hashed' },
        unique: false,
        balancing: true,
        chunkMetadata: [
          { shard: 'a', nChunks: 406 },
          { shard: 'b', nChunks: 405 },
          { shard: 'c', nChunks: 405 }
        ],
        chunks: [
          'too many chunks to print, use verbose if you want to force print'
        ],
        tags: []
      }
    }
  },
  {
    database: {
      _id: 'test',
      primary: 'c',
      partitioned: false,
      version: {
        uuid: new UUID("936d4893-98b7-4189-983d-6da0a60b14f1"),
        timestamp: Timestamp({ t: 1667370651, i: 1 }),
        lastMod: 1
      }
    },
    collections: {}
  },
]

```

Checking the shard distribution:

```
mongos
[direct: mongos] nosql_mid> db.books.getShardDistribution()
Shard a at a/172.31.84.139:27021,172.31.90.164:27021,172.31.93.248:27021
{
  data: '485.21KiB',
  docs: 28512,
  chunks: 406,
  'estimated data per chunk': '1022KiB',
  'estimated docs per chunk': 70
}
---
Shard c at c/172.31.84.139:27023,172.31.90.164:27023,172.31.93.248:27023
{
  data: '398.69KiB',
  docs: 26320,
  chunks: 405,
  'estimated data per chunk': '1008KiB',
  'estimated docs per chunk': 64
}
---
Shard b at b/172.31.84.139:27022,172.31.90.164:27022,172.31.93.248:27022
{
  data: '395.12KiB',
  docs: 26680,
  chunks: 405,
  'estimated data per chunk': '999KiB',
  'estimated docs per chunk': 65
}
---
Totals
{
  data: '1.17GiB',
  docs: 81512,
  chunks: 1216,
  'Shard a': [
    '33.79 % data',
    '34.97 % docs in cluster',
    '104KiB avg obj size on shard'
  ],
  'Shard c': [
    '33.25 % data',
    '32.28 % docs in cluster',
    '15KiB avg obj size on shard'
  ],
  'Shard b': [
    '32.95 % data',
    '32.73 % docs in cluster',
    '15KiB avg obj size on shard'
  ]
}
[direct: mongos] nosql_mid>
```

12. Running Queries on Data

1. Find documents in a given range

Command:

```
db.books.find({ num_likes: { $gte:3000, $lte:3500}}, {_id:1, title:1,
num_likes:1}).explain('executionStats')
```

Query Output:

```
mongos
[direct: mongos] nosql_mid> db.books.find({ num_likes: { $gte:3000, $lte:3500}}, {_id:1, title:1, num_likes:1})
[
  {
    _id: 50,
    title: 'The Best Epic Fantasy (fiction)',
    num_likes: 3078
  },
  {
    _id: 10762, title: 'Best Book Boyfriends', num_likes: 3092 },
  {
    _id: 72680,
    title: 'Gilmore Girls Complete Reading List',
    num_likes: 3491
  },
  {
    _id: 952,
    title: '1801 Books You Must Read Before You Die',
    num_likes: 3386
  },
  {
    _id: 15, title: 'Best Historical Fiction', num_likes: 3178 },
  {
    _id: 1058,
    title: 'Microhistory: Social Histories of Just One Thing',
    num_likes: 3158
  }
]
[direct: mongos] nosql_mid>
```

Execution Stats:

Execution Time: 3900

Documents returned per shard:

- a: 1
- b: 3
- c: 2

Part of Explain command output showing the relevant data:

```

mongoos
x

executionDate: {
  $returned: 0,
  executionTimeMillis: 3900,
  totalKeysExamined: 0,
  totalDocsExamined: 81612,
  executionStages: {
    stage: "PROJECTION_SIMPLE",
    $returned: 0,
    executionTimeMillis: 3900,
    totalKeysExamined: 0,
    totalDocsExamined: 81612,
    totalCpuTimeInMillis: Long("9390"),
    shards: [
      {
        shardId: '0',
        executionSuccess: true,
        $returned: 0,
        executionTimeMillis: 2268,
        totalKeysExamined: 0,
        totalDocsExamined: 26682,
        executionStages: {
          stage: "PROJECTION_FILTER",
          $returned: 0,
          executionTimeMillisEstimate: 2150,
          works: 26682,
          advanced: 2,
          needTime: 26678,
          needYield: 0,
          saveState: 112,
          restoreState: 112,
          iEOF: 1,
          transformBy: { $id: 1, title: 1, num_likes: 1 },
          stage: "SHARDING_FILTER",
          $returned: 0,
          executionTimeMillisEstimate: 2150,
          works: 26682,
          advanced: 2,
          needTime: 26678,
          needYield: 0,
          saveState: 112,
          restoreState: 112,
          iEOF: 1,
          chunkSkips: 0,
          inputStage: {
            stage: "COLSCAN",
            filter: {
              $and: [
                { num_likes: { $lte: 3500 } },
                { num_likes: { $gte: 3000 } } ]
              },
            $returned: 0,
            executionTimeMillisEstimate: 2150,
            works: 26682,
            advanced: 3,
            needTime: 26678,
            needYield: 0,
            saveState: 112,
            restoreState: 112,
            iEOF: 1,
            direction: 'forward',
            docsExamined: 26680
          }
        }
      },
      {
        shardId: '1',
        executionSuccess: true,
        $returned: 1,
        executionTimeMillis: 3236,
        totalKeysExamined: 0,
        totalDocsExamined: 28512,
        executionStages: {
          stage: "PROJECTION_SIMPLE",
          $returned: 2,
          executionTimeMillisEstimate: 2882,
          works: 28514,
          advanced: 1,
          needTime: 28512,
          needYield: 0,
          saveState: 152,
          restoreState: 152,
          iEOF: 1,
          transformBy: { $id: 1, title: 1, num_likes: 1 },
          stage: "SHARDING_FILTER",
          $returned: 1,
          executionTimeMillisEstimate: 2882,
          works: 28514,
          advanced: 1,
          needTime: 28512,
          needYield: 0,
          saveState: 152,
          restoreState: 152,
          iEOF: 1,
          chunkSkips: 0,
          inputStage: {
            stage: "COLSCAN",
            filter: {
              $and: [
                { num_likes: { $lte: 3500 } },
                { num_likes: { $gte: 3000 } } ]
              },
            $returned: 1,
            executionTimeMillisEstimate: 2882,
            works: 28514,
            advanced: 1,
            needTime: 28512,
            needYield: 0,
            saveState: 152,
            restoreState: 152,
            iEOF: 1,
            direction: 'forward',
            docsExamined: 28512
          }
        }
      }
    ],
    $shardTime: 'c',
    executionSuccess: true,
    $returned: 2,
    executionTimeMillis: 3908,
    totalKeysExamined: 0,
    totalDocsExamined: 26320,
    executionStages: {
      stage: "PROJECTION_SIMPLE",
      $returned: 2,
      executionTimeMillisEstimate: 3822,
      works: 26322,
      advanced: 2,
      needTime: 26319,
      needYield: 0,
      saveState: 202,
      restoreState: 202,
      iEOF: 1,
      transformBy: { $id: 1, title: 1, num_likes: 1 },
      inputStage: {
        stage: "SHARDING_FILTER",
        $returned: 2,
        executionTimeMillisEstimate: 3822,
        works: 26322,
        advanced: 2,
        needTime: 26319,
        needYield: 0,
        saveState: 202,
        restoreState: 202,
        iEOF: 1,
        chunkSkips: 0,
        inputStage: {
          stage: "COLSCAN",
          filter: {
            $and: [
              { num_likes: { $lte: 3500 } },
              { num_likes: { $gte: 3000 } } ]
            },
            $returned: 2,
            executionTimeMillisEstimate: 3822,
            works: 26322,
            advanced: 2,
            needTime: 26319,
            needYield: 0,
            saveState: 202,
            restoreState: 202,
            iEOF: 1,
            direction: 'forward',
            docsExamined: 26320
          }
        }
      }
    }
  },
  serverInfo: {
    host: 'ip-172-31-60-79',
    port: 27020,
    version: '3.6.2',
    gitVersion: '96bf7dfe32b77f1f534be7ea394dbf8eeba30e'
  },
  serverParameters: {}
}

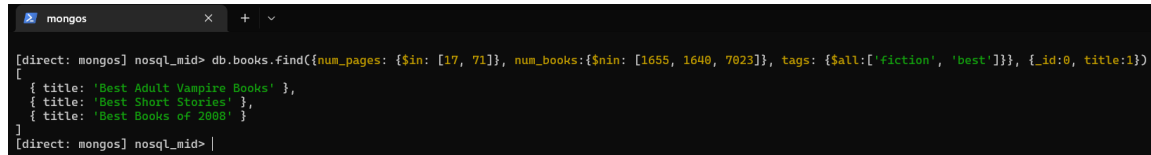
```

2. Query with \$elemMatch with two conditions

Command:

```
db.books.find({num_pages: {$in: [17, 71]}, num_books: {$nin: [1655, 1640, 7023]}, tags:
{$all:['fiction', 'best']}}, {_id:0, title:1}).explain('executionStats')
```

Query Output:



```
mongos
[direct: mongos] nosql_mid> db.books.find({num_pages: {$in: [17, 71]}, num_books: {$nin: [1655, 1640, 7023]}, tags: {$all:['fiction', 'best']}}, {_id:0, title:1})
[
  { title: 'Best Adult Vampire Books' },
  { title: 'Best Short Stories' },
  { title: 'Best Books of 2008' }
]
[direct: mongos] nosql_mid> |
```

Execution Stats:

Execution Time: 314

Documents returned per shard:

a: 0

b: 2

c: 1

Part of Explain command output showing the relevant data:

(next page)

```

mongos
executionStats: {
  nReturned: 3,
  executionTimeMillis: 310,
  totalKeysExamined: 0,
  totalDocsExamined: 81512,
  executionStages: {
    stage: 'PROJECTION_SIMPLE',
    nReturned: 3,
    executionTimeMillis: 310,
    totalKeysExamined: 0,
    totalDocsExamined: 81512,
    totalChillMillis: Long("682"),
    shards: [
      {
        shardName: 'b',
        executionSuccess: true,
        nReturned: 2,
        executionTimeMillis: 100,
        totalKeysExamined: 0,
        totalDocsExamined: 26680,
        executionStages: {
          stage: 'PROJECTION_SIMPLE',
          nReturned: 2,
          executionTimeMillisEstimate: 103,
          works: 26682,
          advanced: 2,
          needTime: 26079,
          needYield: 0,
          saveState: 20,
          restoreState: 20,
          isEOF: 1,
          transformBy: { _id: 0, title: 1 },
          inputStage: {
            stage: 'SHARDING_FILTER',
            nReturned: 2,
            executionTimeMillisEstimate: 103,
            works: 26682,
            advanced: 2,
            needTime: 26079,
            needYield: 0,
            saveState: 20,
            restoreState: 20,
            isEOF: 1,
            chunkSkips: 0,
            inputStage: {
              stage: 'COLSCAN',
              filter: {
                $and: [
                  { tags: { $eq: 'fiction' } },
                  { tags: { $eq: 'best' } },
                  { num_pages: { $in: [ 17, 71 ] } },
                  { num_books: { $not: { $in: [ 1640, 1655, 7023 ] } } }
                ]
              }
            },
            nReturned: 2,
            executionTimeMillisEstimate: 103,
            works: 26682,
            advanced: 2,
            needTime: 26079,
            needYield: 0,
            saveState: 20,
            restoreState: 20,
            isEOF: 1,
            direction: 'forward',
            docsExamined: 26680
          }
        ]
      },
      {
        shardName: 'a',
        executionSuccess: true,
        nReturned: 0,
        executionTimeMillis: 220,
        totalKeysExamined: 0,
        totalDocsExamined: 28512,
        executionStages: {
          stage: 'PROJECTION_SIMPLE',
          nReturned: 0,
          executionTimeMillisEstimate: 176,
          works: 28514,
          advanced: 0,
          needTime: 28513,
          needYield: 0,
          saveState: 31,
          restoreState: 31,
          isEOF: 1,
          transformBy: { _id: 0, title: 1 },
          inputStage: {
            stage: 'SHARDING_FILTER',
            nReturned: 0,
            executionTimeMillisEstimate: 176,
            works: 28516,
            advanced: 0,
            needTime: 28513,
            needYield: 0,
            saveState: 31,
            restoreState: 31,
            isEOF: 1,
            chunkSkips: 0,
            inputStage: {
              stage: 'COLSCAN',
              filter: {
                $and: [
                  { tags: { $eq: 'fiction' } },
                  { tags: { $eq: 'best' } },
                  { num_pages: { $in: [ 17, 71 ] } },
                  { num_books: { $not: { $in: [ 1640, 1655, 7023 ] } } }
                ]
              }
            },
            nReturned: 0,
            executionTimeMillisEstimate: 176,
            works: 28514,
            advanced: 0,
            needTime: 28513,
            needYield: 0,
            saveState: 31,
            restoreState: 31,
            isEOF: 1,
            direction: 'forward',
            docsExamined: 28512
          }
        ]
      }
    ]
  },
  shardName: 'c',
  executionSuccess: true,
  nReturned: 1,
  executionTimeMillis: 312,
  totalKeysExamined: 0,
  totalDocsExamined: 26320,
  executionStages: {
    stage: 'PROJECTION_SIMPLE',
    nReturned: 1,
    executionTimeMillisEstimate: 238,
    works: 26322,
    advanced: 1,
    needTime: 26320,
    needYield: 0,
    saveState: 31,
    restoreState: 31,
    isEOF: 1,
    transformBy: { _id: 0, title: 1 },
    inputStage: {
      stage: 'SHARDING_FILTER',
      nReturned: 1,
      executionTimeMillisEstimate: 238,
      works: 26322,
      advanced: 1,
      needTime: 26320,
      needYield: 0,
      saveState: 31,
      restoreState: 31,
      isEOF: 1,
      chunkSkips: 0,
      inputStage: {
        stage: 'COLSCAN',
        filter: {
          $and: [
            { tags: { $eq: 'fiction' } },
            { tags: { $eq: 'best' } },
            { num_pages: { $in: [ 17, 71 ] } },
            { num_books: { $not: { $in: [ 1640, 1655, 7023 ] } } }
          ]
        }
      },
      nReturned: 1,
      executionTimeMillisEstimate: 238,
      works: 26322,
      advanced: 1,
      needTime: 26320,
      needYield: 0,
      saveState: 31,
      restoreState: 31,
      isEOF: 1,
      direction: 'forward',
      docsExamined: 26320
    }
  }
}

```


3. Query with \$in, \$nin, or \$all

Command:

```
db.books.find({books: {"$elemMatch": {"position.ranking": {$gt: 2000, $lt: 2050},  
author_id: {$in: ["3618", "7628"]}}}}, {_id: 0, title: 1}).explain('executionStats')
```

Query Output:

```
mongos
[direct: mongos] nosql_mid> db.books.find({books: {"$elemMatch": {"position.ranking": {$gt: 2000, $lt: 2050}, author_id: {$in: ["3618", "7628"]}}}}, {_id: 0, title: 1})
[
  { title: 'What Book Would You Like To Live In?' },
  { title: 'Best Books of the 21st Century' },
  { title: 'Books that Changed the Way You View Life' },
  { title: 'Books I Became So Obsessed With I Stopped Everything Else in My Life to Finish.' }
]
[direct: mongos] nosql_mid> |
```

Execution Stats:

Execution Time: 1218

Documents returned per shard:

a: 2

b: 1

c: 1

Part of Explain command output showing the relevant data:

(next page)

```

executionStats: {
  returned: 0,
  executionTimeMillis: 1218,
  totalTimeExamined: 0,
  totalDocsExamined: 81512,
  executionStages: {
    stage: 'QUERY_INITIALIZE',
    nReturned: 0,
    executionTimeMillis: 1218,
    totalTimeExamined: 0,
    totalDocsExamined: 81512,
    totalChildMillis: Long("3951"),
    shards: [
      {
        shardName: 'c',
        executionSuccess: true,
        nReturned: 1,
        executionTimeMillis: 1197,
        totalTimeExamined: 0,
        totalDocsExamined: 26680,
        executionStages: {
          stage: 'QUERY_INITIALIZE',
          nReturned: 1,
          executionTimeMillisEstimate: 1125,
          works: 26682,
          advanced: 1,
          needTime: 26680,
          needYield: 0,
          saveState: 60,
          restoreState: 60,
          isEOF: 1,
          transformBy: { _id: 0, title: 1 },
          inputStage: {
            stage: 'SHARDING_FILTER',
            nReturned: 1,
            executionTimeMillisEstimate: 1125,
            works: 26682,
            advanced: 1,
            needTime: 26680,
            needYield: 0,
            saveState: 60,
            restoreState: 60,
            isEOF: 1,
            chunkSkips: 0,
            inputStage: {
              stage: 'COLSCAN',
              filter: {
                books: {
                  $elemMatch: {
                    'name': {
                      { position: ranking: { '$lt': 2050 } },
                      { position: ranking: { '$gt': 2000 } },
                      { author_id: { '$in': [ '3010', '3020' ] } }
                    }
                  }
                }
              }
            }
          },
          nReturned: 1,
          executionTimeMillisEstimate: 1125,
          works: 26682,
          advanced: 1,
          needTime: 26680,
          needYield: 0,
          saveState: 60,
          restoreState: 60,
          isEOF: 1,
          direction: 'forward',
          docsExamined: 26680
        }
      }
    ]
  },
  shardName: 'c',
  executionSuccess: true,
  nReturned: 1,
  executionTimeMillis: 1190,
  totalTimeExamined: 0,
  totalDocsExamined: 26320,
  executionStages: {
    stage: 'QUERY_INITIALIZE',
    nReturned: 1,
    executionTimeMillisEstimate: 1186,
    works: 26322,
    advanced: 1,
    needTime: 26320,
    needYield: 0,
    saveState: 62,
    restoreState: 62,
    isEOF: 1,
    transformBy: { _id: 0, title: 1 },
    inputStage: {
      stage: 'SHARDING_FILTER',
      nReturned: 1,
      executionTimeMillisEstimate: 1185,
      works: 26322,
      advanced: 1,
      needTime: 26320,
      needYield: 0,
      saveState: 62,
      restoreState: 62,
      isEOF: 1,
      chunkSkips: 0,
      inputStage: {
        stage: 'COLSCAN',
        filter: {
          books: {
            $elemMatch: {
              'name': {
                { position: ranking: { '$lt': 2050 } },
                { position: ranking: { '$gt': 2000 } },
                { author_id: { '$in': [ '3010', '3020' ] } }
              }
            }
          }
        }
      }
    },
    nReturned: 1,
    executionTimeMillisEstimate: 1185,
    works: 26322,
    advanced: 1,
    needTime: 26320,
    needYield: 0,
    saveState: 62,
    restoreState: 62,
    isEOF: 1,
    direction: 'forward',
    docsExamined: 26320
  }
},
  shardName: 'a',
  executionSuccess: true,
  nReturned: 2,
  executionTimeMillis: 1216,
  totalTimeExamined: 0,
  totalDocsExamined: 28512,
  executionStages: {
    stage: 'QUERY_INITIALIZE',
    nReturned: 2,
    executionTimeMillisEstimate: 1207,
    works: 28514,
    advanced: 2,
    needTime: 28511,
    needYield: 0,
    saveState: 66,
    restoreState: 66,
    isEOF: 1,
    transformBy: { _id: 0, title: 1 },
    inputStage: {
      stage: 'SHARDING_FILTER',
      nReturned: 2,
      executionTimeMillisEstimate: 1207,
      works: 28514,
      advanced: 2,
      needTime: 28511,
      needYield: 0,
      saveState: 66,
      restoreState: 66,
      isEOF: 1,
      chunkSkips: 0,
      inputStage: {
        stage: 'COLSCAN',
        filter: {
          books: {
            $elemMatch: {
              'name': {
                { position: ranking: { '$lt': 2050 } },
                { position: ranking: { '$gt': 2000 } },
                { author_id: { '$in': [ '3010', '3020' ] } }
              }
            }
          }
        }
      }
    },
    nReturned: 2,
    executionTimeMillisEstimate: 1207,
    works: 28514,
    advanced: 2,
    needTime: 28511,
    needYield: 0,
    saveState: 66,
    restoreState: 66,
    isEOF: 1,
    direction: 'forward',
    docsExamined: 28512
  }
}
}

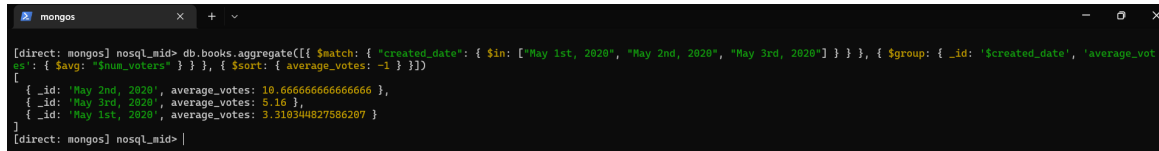
```

4. Query involving aggregate()

Command:

```
db.books.aggregate([ { $match: { "created_date": { $in: ["May 1st, 2020", "May 2nd, 2020", "May 3rd, 2020"] } } }, { $group: { _id: '$created_date', 'average_votes': { $avg: "$num_voters" } } }, { $sort: { average_votes: -1 } } ] ).explain('executionStats')
```

Query Output:



```
[direct: mongos] nosql_mid> db.books.aggregate([ { $match: { "created_date": { $in: ["May 1st, 2020", "May 2nd, 2020", "May 3rd, 2020"] } } }, { $group: { _id: '$created_date', 'average_votes': { $avg: "$num_voters" } } }, { $sort: { average_votes: -1 } } ] )
[
  { _id: 'May 2nd, 2020', average_votes: 18.666666666666666 },
  { _id: 'May 3rd, 2020', average_votes: 5.16 },
  { _id: 'May 1st, 2020', average_votes: 3.318344827586207 }
]
[direct: mongos] nosql_mid>
```

Execution Stats:

Execution Time:

a: 175 ms

b: 97 ms

c: 158 ms

Documents returned per shard:

a: 3

b: 3

c: 3

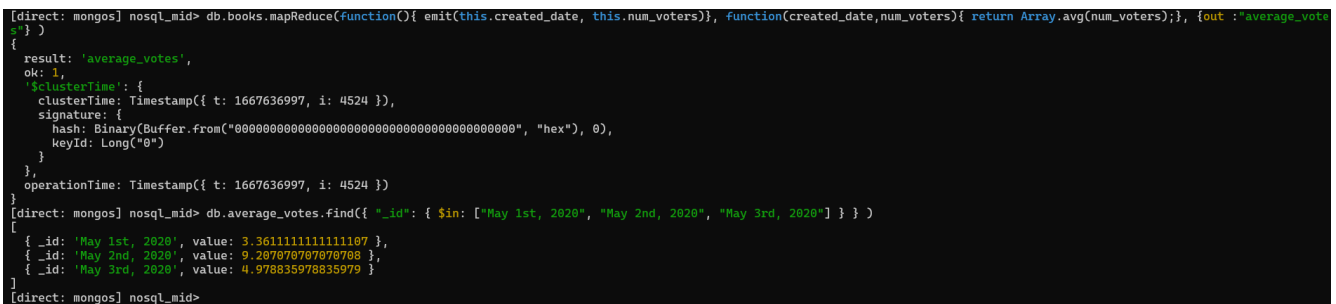
5. Query involving mapReduce()

Command:

```
var m_start = new Date(); db.books.mapReduce(function() { emit(this.created_date, this.num_voters); }, function(created_date, num_voters) { return Array.avg(num_voters); }, { out: "average_votes" }); var m_end = new Date(); execution_10 = m_end - m_start
```

```
db.average_votes.find( { "_id": { $in: ["May 1st, 2020", "May 2nd, 2020", "May 3rd, 2020"] } } )
```

Query Output:



```
[direct: mongos] nosql_mid> db.books.mapReduce(function() { emit(this.created_date, this.num_voters); }, function(created_date, num_voters) { return Array.avg(num_voters); }, { out: "average_votes" })
{
  result: 'average_votes',
  ok: 1,
  '$clusterTime': {
    clusterTime: Timestamp({ t: 1667636997, i: 4524 }),
    signature: {
      hash: Binary(Buffer.from("00000000000000000000000000000000", "hex"), 0),
      keyId: Long("0")
    }
  },
  operationTime: Timestamp({ t: 1667636997, i: 4524 })
}
[direct: mongos] nosql_mid> db.average_votes.find( { "_id": { $in: ["May 1st, 2020", "May 2nd, 2020", "May 3rd, 2020"] } } )
[
  { _id: 'May 1st, 2020', value: 3.361111111111111 },
  { _id: 'May 2nd, 2020', value: 9.287878787878788 },
  { _id: 'May 3rd, 2020', value: 4.978835978835979 }
]
[direct: mongos] nosql_mid>
```

Execution Stats:

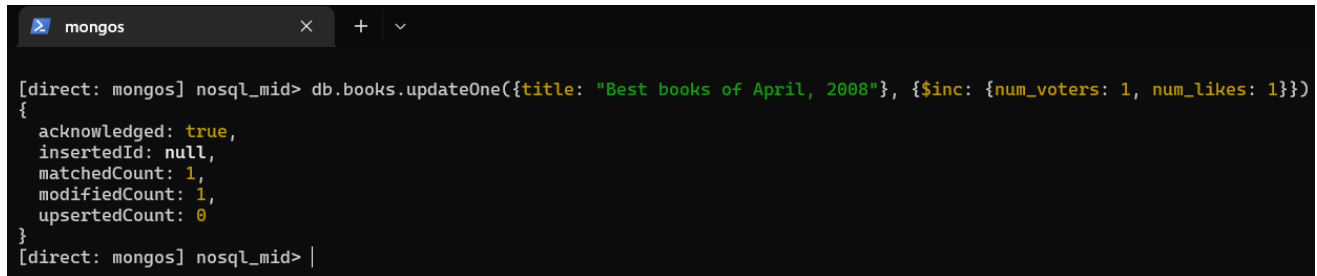
Execution time: 2343 ms

6. Update query

Command:

```
var m_start = new Date(); db.books.updateOne({title: "Best books of April, 2008"}, {$inc: {num_voters: 1, num_likes: 1}}); var m_end = new Date(); execution_10 = m_end - m_start
```

Query Output:

A screenshot of a MongoDB shell window titled 'mongos'. The command executed is `db.books.updateOne({title: "Best books of April, 2008"}, {$inc: {num_voters: 1, num_likes: 1}})`. The output is a JSON object: `{ acknowledged: true, insertedId: null, matchedCount: 1, modifiedCount: 1, upsertedCount: 0 }`. The prompt is `[direct: mongos] nosql_mid> |`.

```
[direct: mongos] nosql_mid> db.books.updateOne({title: "Best books of April, 2008"}, {$inc: {num_voters: 1, num_likes: 1}})
{
  acknowledged: true,
  insertedId: null,
  matchedCount: 1,
  modifiedCount: 1,
  upsertedCount: 0
}
[direct: mongos] nosql_mid> |
```

Execution Stats:

Execution time: 12ms

13. Replication of Shards

// rs.status image of shards

Replication Status of 3 Shards: shard a, shard b, & shard c:
(next three pages)

```

a [direct: primary] test> rs.status()
{
  set: 'a',
  date: ISODate("2022-11-05T06:38:15.319Z"),
  myState: 1,
  term: Long("14"),
  syncSourceHost: '',
  syncSourceId: -1,
  heartbeatIntervalMillis: Long("2000"),
  majorityVoteCount: 2,
  writeMajorityCount: 2,
  votingMembersCount: 3,
  writableVotingMembersCount: 3,
  optimes: {
    lastCommittedOpTime: { ts: Timestamp({ t: 1667630287, i: 1 }), t: Long("14") },
    lastCommittedWallTime: ISODate("2022-11-05T06:38:07.546Z"),
    readConcernMajorityOpTime: { ts: Timestamp({ t: 1667630287, i: 1 }), t: Long("14") },
    appliedOpTime: { ts: Timestamp({ t: 1667630287, i: 1 }), t: Long("14") },
    durableOpTime: { ts: Timestamp({ t: 1667630287, i: 1 }), t: Long("14") },
    lastAppliedWallTime: ISODate("2022-11-05T06:38:07.546Z"),
    lastDurableWallTime: ISODate("2022-11-05T06:38:07.546Z")
  },
  lastStableRecoveryTimestamp: Timestamp({ t: 1667630237, i: 1 }),
  electionCandidateMetrics: {
    lastElectionReason: 'priorityTakeover',
    lastElectionDate: ISODate("2022-11-05T06:16:37.482Z"),
    electionTerm: Long("14"),
    lastCommittedOpTimeAtElection: { ts: Timestamp({ t: 1667628996, i: 1 }), t: Long("13") },
    lastSeenOpTimeAtElection: { ts: Timestamp({ t: 1667628996, i: 1 }), t: Long("13") },
    numVotesNeeded: 2,
    priorityAtElection: 4,
    electionTimeoutMillis: Long("10000"),
    priorPrimaryMemberId: 2,
    numCatchUpOps: Long("0"),
    newTermStartDate: ISODate("2022-11-05T06:16:37.495Z"),
    wMajorityWriteAvailabilityDate: ISODate("2022-11-05T06:16:38.494Z")
  },
  electionParticipantMetrics: {
    votedForCandidate: true,
    electionTerm: Long("13"),
    lastVoteDate: ISODate("2022-11-05T06:16:26.368Z"),
    electionCandidateMemberId: 2,
    voteReason: '',
    lastAppliedOpTimeAtElection: { ts: Timestamp({ t: 1667540934, i: 2 }), t: Long("11") },
    maxAppliedOpTimeInSet: { ts: Timestamp({ t: 1667540934, i: 2 }), t: Long("11") },
    priorityAtElection: 4
  },
  members: [
    {
      _id: 0,
      name: '172.31.93.248:27021',
      health: 1,
      state: 1,
      stateStr: 'PRIMARY',
      uptime: 1318,
      optime: { ts: Timestamp({ t: 1667630287, i: 1 }), t: Long("14") },
      optimeDate: ISODate("2022-11-05T06:38:07.000Z"),
      lastAppliedWallTime: ISODate("2022-11-05T06:38:07.546Z"),
      lastDurableWallTime: ISODate("2022-11-05T06:38:07.546Z"),
      syncSourceHost: '',
      syncSourceId: -1,
      infoMessage: '',
      electionTime: Timestamp({ t: 1667628997, i: 1 }),
      electionDate: ISODate("2022-11-05T06:16:37.000Z"),
      configVersion: 2,
      configTerm: 14,
      self: true,
      lastHeartbeatMessage: ''
    },
    {
      _id: 1,
      name: '172.31.84.139:27021',
      health: 1,
      state: 2,
      stateStr: 'SECONDARY',
      uptime: 1309,
      optime: { ts: Timestamp({ t: 1667630287, i: 1 }), t: Long("14") },
      optimeDurable: { ts: Timestamp({ t: 1667630287, i: 1 }), t: Long("14") },
      optimeDate: ISODate("2022-11-05T06:38:07.000Z"),
      optimeDurableDate: ISODate("2022-11-05T06:38:07.000Z"),
      lastAppliedWallTime: ISODate("2022-11-05T06:38:07.546Z"),
      lastDurableWallTime: ISODate("2022-11-05T06:38:07.546Z"),
      lastHeartbeat: ISODate("2022-11-05T06:38:14.025Z"),
      lastHeartbeatRecv: ISODate("2022-11-05T06:38:14.024Z"),
      pingMs: Long("0"),
      lastHeartbeatMessage: '',
      syncSourceHost: '172.31.90.164:27021',
      syncSourceId: 2,
      infoMessage: '',
      configVersion: 2,
      configTerm: 14
    },
    {
      _id: 2,
      name: '172.31.90.164:27021',
      health: 1,
      state: 2,
      stateStr: 'SECONDARY',
      uptime: 1315,
      optime: { ts: Timestamp({ t: 1667630287, i: 1 }), t: Long("14") },
      optimeDurable: { ts: Timestamp({ t: 1667630287, i: 1 }), t: Long("14") },
      optimeDate: ISODate("2022-11-05T06:38:07.000Z"),
      optimeDurableDate: ISODate("2022-11-05T06:38:07.000Z"),
      lastAppliedWallTime: ISODate("2022-11-05T06:38:07.546Z"),
      lastDurableWallTime: ISODate("2022-11-05T06:38:07.546Z"),
      lastHeartbeat: ISODate("2022-11-05T06:38:14.026Z"),
      lastHeartbeatRecv: ISODate("2022-11-05T06:38:15.163Z"),
      pingMs: Long("0"),
      lastHeartbeatMessage: '',
      syncSourceHost: '172.31.93.248:27021',
      syncSourceId: 0,
      infoMessage: '',
      configVersion: 2,
      configTerm: 14
    }
  ],
  ok: 1,
  lastCommittedOpTime: Timestamp({ t: 1667630287, i: 1 }),
  '$clusterTime': {
    clusterTime: Timestamp({ t: 1667630289, i: 1 }),
    signature: {
      hash: Binary(Buffer.from("00000000000000000000000000000000", "hex"), 0),
      keyId: Long("0")
    }
  }
}

```

```

b [direct: primary] test> rs.status()
{
  set: 'b',
  date: ISODate("2022-11-05T06:38:31.641Z"),
  myState: 1,
  term: Long("13"),
  syncSourceHost: '',
  syncSourceId: -1,
  heartbeatIntervalMillis: Long("2000"),
  majorityVoteCount: 2,
  writeMajorityCount: 2,
  votingMembersCount: 3,
  writableVotingMembersCount: 3,
  optimes: {
    lastCommittedOpTime: { ts: Timestamp({ t: 1667630302, i: 1 }), t: Long("13") },
    lastCommittedWallTime: ISODate("2022-11-05T06:38:22.879Z"),
    readConcernMajorityOpTime: { ts: Timestamp({ t: 1667630302, i: 1 }), t: Long("13") },
    appliedOpTime: { ts: Timestamp({ t: 1667630302, i: 1 }), t: Long("13") },
    durableOpTime: { ts: Timestamp({ t: 1667630302, i: 1 }), t: Long("13") },
    lastAppliedWallTime: ISODate("2022-11-05T06:38:22.879Z"),
    lastDurableWallTime: ISODate("2022-11-05T06:38:22.879Z")
  },
  lastStableRecoveryTimestamp: Timestamp({ t: 1667630262, i: 1 }),
  electionCandidateMetrics: {
    lastElectionReason: 'priorityTakeover',
    lastElectionDate: ISODate("2022-11-05T06:18:02.783Z"),
    electionTerm: Long("13"),
    lastCommittedOpTimeAtElection: { ts: Timestamp({ t: 1667629081, i: 1 }), t: Long("12") },
    lastSeenOpTimeAtElection: { ts: Timestamp({ t: 1667629081, i: 1 }), t: Long("12") },
    numVotesNeeded: 2,
    priorityAtElection: 4,
    electionTimeoutMillis: Long("10000"),
    priorPrimaryMemberId: 0,
    numCatchUpOps: Long("0"),
    newTermStartDate: ISODate("2022-11-05T06:18:02.819Z"),
    wMajorityWriteAvailabilityDate: ISODate("2022-11-05T06:18:03.873Z")
  },
  electionParticipantMetrics: {
    votedForCandidate: true,
    electionTerm: Long("12"),
    lastVoteDate: ISODate("2022-11-05T06:17:51.649Z"),
    electionCandidateMemberId: 0,
    voteReason: '',
    lastAppliedOpTimeAtElection: { ts: Timestamp({ t: 1667540932, i: 1 }), t: Long("10") },
    maxAppliedOpTimeInSet: { ts: Timestamp({ t: 1667540932, i: 1 }), t: Long("10") },
    priorityAtElection: 4
  },
  members: [
    {
      _id: 0,
      name: '172.31.93.248:27022',
      health: 1,
      state: 2,
      stateStr: 'SECONDARY',
      uptime: 1245,
      optime: { ts: Timestamp({ t: 1667630302, i: 1 }), t: Long("13") },
      optimeDurable: { ts: Timestamp({ t: 1667630302, i: 1 }), t: Long("13") },
      optimeDate: ISODate("2022-11-05T06:38:22.000Z"),
      optimeDurableDate: ISODate("2022-11-05T06:38:22.000Z"),
      lastAppliedWallTime: ISODate("2022-11-05T06:38:22.879Z"),
      lastDurableWallTime: ISODate("2022-11-05T06:38:22.879Z"),
      lastHeartbeat: ISODate("2022-11-05T06:38:31.370Z"),
      lastHeartbeatRecv: ISODate("2022-11-05T06:38:30.623Z"),
      pingMs: Long("0"),
      lastHeartbeatMessage: '',
      syncSourceHost: '172.31.84.139:27022',
      syncSourceId: 1,
      infoMessage: '',
      configVersion: 2,
      configTerm: 13
    },
    {
      _id: 1,
      name: '172.31.84.139:27022',
      health: 1,
      state: 1,
      stateStr: 'PRIMARY',
      uptime: 1247,
      optime: { ts: Timestamp({ t: 1667630302, i: 1 }), t: Long("13") },
      optimeDate: ISODate("2022-11-05T06:38:22.000Z"),
      lastAppliedWallTime: ISODate("2022-11-05T06:38:22.879Z"),
      lastDurableWallTime: ISODate("2022-11-05T06:38:22.879Z"),
      syncSourceHost: '',
      syncSourceId: -1,
      infoMessage: '',
      electionTime: Timestamp({ t: 1667629082, i: 1 }),
      electionDate: ISODate("2022-11-05T06:18:02.000Z"),
      configVersion: 2,
      configTerm: 13,
      self: true,
      lastHeartbeatMessage: ''
    },
    {
      _id: 2,
      name: '172.31.90.164:27022',
      health: 1,
      state: 2,
      stateStr: 'SECONDARY',
      uptime: 1245,
      optime: { ts: Timestamp({ t: 1667630302, i: 1 }), t: Long("13") },
      optimeDurable: { ts: Timestamp({ t: 1667630302, i: 1 }), t: Long("13") },
      optimeDate: ISODate("2022-11-05T06:38:22.000Z"),
      optimeDurableDate: ISODate("2022-11-05T06:38:22.000Z"),
      lastAppliedWallTime: ISODate("2022-11-05T06:38:22.879Z"),
      lastDurableWallTime: ISODate("2022-11-05T06:38:22.879Z"),
      lastHeartbeat: ISODate("2022-11-05T06:38:31.369Z"),
      lastHeartbeatRecv: ISODate("2022-11-05T06:38:31.413Z"),
      pingMs: Long("0"),
      lastHeartbeatMessage: '',
      syncSourceHost: '172.31.93.248:27022',
      syncSourceId: 0,
      infoMessage: '',
      configVersion: 2,
      configTerm: 13
    }
  ],
  ok: 1,
  lastCommittedOpTime: Timestamp({ t: 1667630302, i: 1 }),
  '$clusterTime': {
    clusterTime: Timestamp({ t: 1667630309, i: 1 }),
    signature: {
      hash: Binary(Buffer.from("00000000000000000000000000000000", "hex"), 0),
      keyId: Long("0")
    }
  }
}

```

```

c [direct: primary] test> rs.status()
{
  set: 'c',
  date: ISODate("2022-11-05T06:38:40.783Z"),
  myState: 1,
  term: Long("17"),
  syncSourceHost: '',
  syncSourceId: -1,
  heartbeatIntervalMillis: Long("2000"),
  majorityVoteCount: 2,
  writeMajorityCount: 2,
  votingMembersCount: 3,
  writableVotingMembersCount: 3,
  optimes: {
    lastCommittedOptime: { ts: Timestamp({ t: 1667630315, i: 1 }), t: Long("17") },
    lastCommittedWallTime: ISODate("2022-11-05T06:38:35.060Z"),
    readConcernMajorityOptime: { ts: Timestamp({ t: 1667630315, i: 1 }), t: Long("17") },
    appliedOptime: { ts: Timestamp({ t: 1667630315, i: 1 }), t: Long("17") },
    durableOptime: { ts: Timestamp({ t: 1667630315, i: 1 }), t: Long("17") },
    lastAppliedWallTime: ISODate("2022-11-05T06:38:35.060Z"),
    lastDurableWallTime: ISODate("2022-11-05T06:38:35.060Z")
  },
  lastStableRecoveryTimestamp: Timestamp({ t: 1667630266, i: 2 }),
  electionCandidateMetrics: {
    lastElectionReason: 'priorityTakeover',
    lastElectionDate: ISODate("2022-11-05T06:18:24.977Z"),
    electionTerm: Long("17"),
    lastCommittedOptimeAtElection: { ts: Timestamp({ t: 1667629103, i: 1 }), t: Long("16") },
    lastSeenOptimeAtElection: { ts: Timestamp({ t: 1667629103, i: 1 }), t: Long("16") },
    numVotesNeeded: 2,
    priorityAtElection: 4,
    electionTimeoutMillis: Long("10000"),
    priorPrimaryMemberId: 0,
    numCatchUpOps: Long("0"),
    newTermStartDate: ISODate("2022-11-05T06:18:24.990Z"),
    wMajorityWriteAvailabilityDate: ISODate("2022-11-05T06:18:25.992Z")
  },
  electionParticipantMetrics: {
    votedForCandidate: true,
    electionTerm: Long("16"),
    lastVoteDate: ISODate("2022-11-05T06:18:13.542Z"),
    electionCandidateMemberId: 0,
    voteReason: '',
    lastAppliedOptimeAtElection: { ts: Timestamp({ t: 1667540927, i: 1 }), t: Long("14") },
    maxAppliedOptimeInSet: { ts: Timestamp({ t: 1667540927, i: 1 }), t: Long("14") },
    priorityAtElection: 4
  },
  members: [
    {
      _id: 0,
      name: '172.31.93.248:27023',
      health: 1,
      state: 2,
      stateStr: 'SECONDARY',
      uptime: 1237,
      optime: { ts: Timestamp({ t: 1667630315, i: 1 }), t: Long("17") },
      optimeDurable: { ts: Timestamp({ t: 1667630315, i: 1 }), t: Long("17") },
      optimeDate: ISODate("2022-11-05T06:38:35.000Z"),
      optimeDurableDate: ISODate("2022-11-05T06:38:35.000Z"),
      lastAppliedWallTime: ISODate("2022-11-05T06:38:35.060Z"),
      lastDurableWallTime: ISODate("2022-11-05T06:38:35.060Z"),
      lastHeartbeat: ISODate("2022-11-05T06:38:39.559Z"),
      lastHeartbeatRecv: ISODate("2022-11-05T06:38:40.662Z"),
      pingMs: Long("0"),
      lastHeartbeatMessage: '',
      syncSourceHost: '172.31.90.164:27023',
      syncSourceId: 2,
      infoMessage: '',
      configVersion: 2,
      configTerm: 17
    },
    {
      _id: 1,
      name: '172.31.84.139:27023',
      health: 1,
      state: 2,
      stateStr: 'SECONDARY',
      uptime: 1234,
      optime: { ts: Timestamp({ t: 1667630315, i: 1 }), t: Long("17") },
      optimeDurable: { ts: Timestamp({ t: 1667630315, i: 1 }), t: Long("17") },
      optimeDate: ISODate("2022-11-05T06:38:35.000Z"),
      optimeDurableDate: ISODate("2022-11-05T06:38:35.000Z"),
      lastAppliedWallTime: ISODate("2022-11-05T06:38:35.060Z"),
      lastDurableWallTime: ISODate("2022-11-05T06:38:35.060Z"),
      lastHeartbeat: ISODate("2022-11-05T06:38:39.559Z"),
      lastHeartbeatRecv: ISODate("2022-11-05T06:38:39.558Z"),
      pingMs: Long("0"),
      lastHeartbeatMessage: '',
      syncSourceHost: '172.31.90.164:27023',
      syncSourceId: 2,
      infoMessage: '',
      configVersion: 2,
      configTerm: 17
    },
    {
      _id: 2,
      name: '172.31.90.164:27023',
      health: 1,
      state: 1,
      stateStr: 'PRIMARY',
      uptime: 1248,
      optime: { ts: Timestamp({ t: 1667630315, i: 1 }), t: Long("17") },
      optimeDate: ISODate("2022-11-05T06:38:35.000Z"),
      lastAppliedWallTime: ISODate("2022-11-05T06:38:35.060Z"),
      lastDurableWallTime: ISODate("2022-11-05T06:38:35.060Z"),
      syncSourceHost: '',
      syncSourceId: -1,
      infoMessage: '',
      electionTime: Timestamp({ t: 1667629104, i: 1 }),
      electionDate: ISODate("2022-11-05T06:18:24.000Z"),
      configVersion: 2,
      configTerm: 17,
      self: true,
      lastHeartbeatMessage: ''
    }
  ],
  ok: 1,
  lastCommittedOptime: Timestamp({ t: 1667630315, i: 1 }),
  $clusterTime: {
    clusterTime: Timestamp({ t: 1667630319, i: 1 }),
    signature: {
      hash: Binary(Buffer.from("00000000000000000000000000000000", "hex"), 0),
      keyId: Long("0")
    }
  }
}

```

14. List of Hosts

Please see the [Deployment Summary](#) in part 1 for a detailed explanation.

Node Number	Instance/ Node Name	Port	Service	IP Address (Private)
Node 1	mongos	27020	Shard Controller	172.31.80.79
Node 2	config1	27019	Config Server Primary	172.31.90.97
Node 3	config2	27019	Config Server Secondary 1	172.31.89.25
Node 4	config3	27019	Config Server Secondary 2	172.31.86.144
Node 5	shardNode1	27021	Shard a Primary	172.31.93.248
Node 5	shardNode1	27022	Shard b Secondary 1	172.31.93.248
Node 5	shardNode1	27023	Shard c Secondary 1	172.31.93.248
Node 6	shardNode2	27021	Shard a Secondary 1	172.31.84.139
Node 6	shardNode2	27022	Shard b Primary	172.31.84.139
Node 6	shardNode2	27023	Shard c Secondary 2	172.31.84.139
Node 7	shardNode3	27021	Shard a Secondary 2	172.31.90.164
Node 7	shardNode3	27022	Shard b Secondary 2	172.31.90.164
Node 7	shardNode3	27023	Shard c Primary	172.31.90.164